

CMPS 405: OPERATING SYSTEMS

Synchronization

This material is based on the operating system books by Silberschatz and Stallings and other Linux and system programming books.

Objectives

❖ Motivations

- Concurrent Programming
 - **Producer/Consumer Problem (PCP)**
- Synchronization Terminology
- Cooperative Processes Interaction Structure
- The Concept of a Lock
- Conditions for a Valid Solution

Objectives

- ❖ Implicit reentrant locks using **synchronized** keyword
 - Solutions of PCP
 - using **synchronized** with **yield** method of **Thread**
 - using **synchronized** with **wait/notify** and **wait/notifyAll** methods of **Object**

- ❖ Explicit reentrant lock in **java.util.concurrent.locks**
 - Solution of PCP
 - using **ReentrantLock** with **Condition** classes of the package **java.util.concurrent.locks**


Objectives

❖ Synchronization Semaphores

- The concept of semaphores
- The **Semaphore** class in **java.util.concurrent.locks** package
 - Solutions to the Dining Philosophers Problem (DPP) using semaphores.
 - Solution of the Producer/Consumer Problem using semaphores.

Concurrent Programming

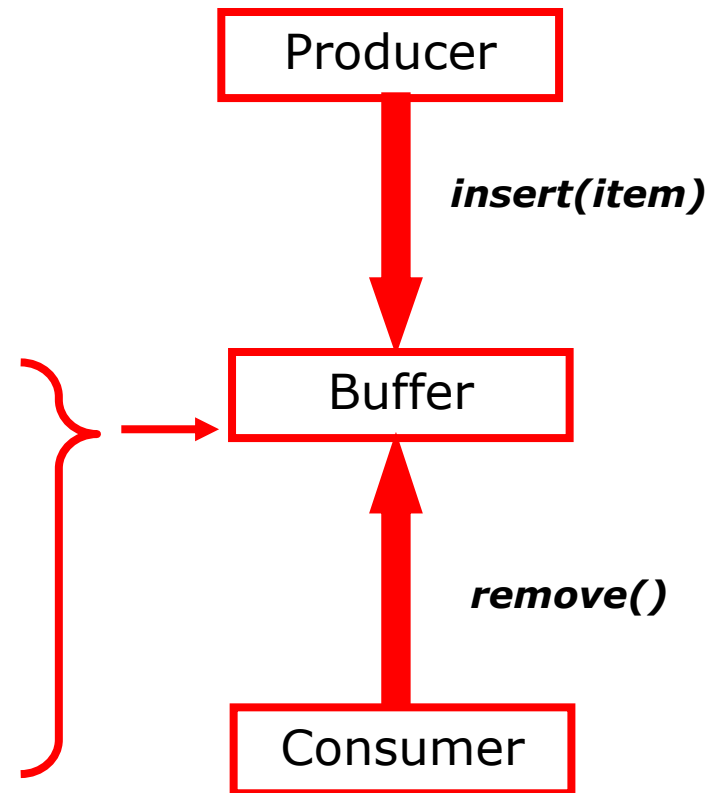
- ❖ In sequential programming: all computational tasks are executed in sequence, one after the other.
- ❖ In concurrent programming: multiple computational tasks are executed simultaneously, at the same time.
- ❖ Implementation of concurrent tasks:
 - as separate programs
 - as a set of processes or threads created by a single program
- ❖ Execution of concurrent tasks:
 - on a single processor
 - ***Multithreaded programming***
 - on several processors in close proximity
 - ***Parallel computing***
 - on several processors distributed across a network
 - ***Distributed computing***



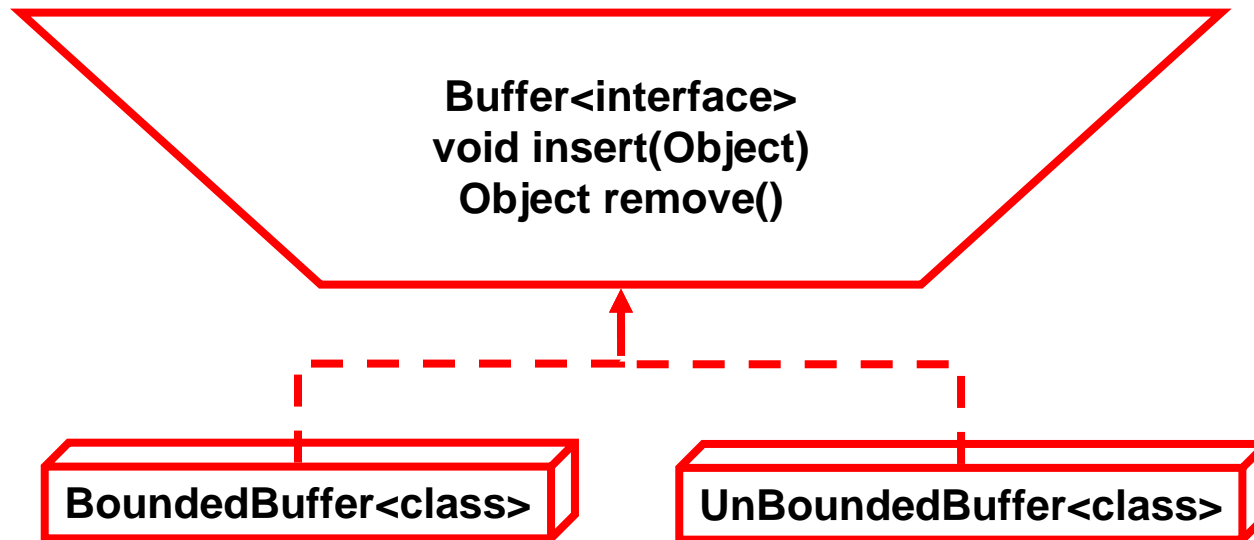
Producer/Consumer Problem

❖ Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process.

- *unbounded-buffer* places no practical limit on the size of the buffer.
- *bounded-buffer* assumes that there is a fixed buffer size.



Modeling The Shared Buffer




```

public interface Buffer
{
    // producers call this method
    public abstract void insert(Object item);

    // consumers call this method
    public abstract Object remove();
}
  
```

Bounded Buffer Implementation

```
public class BoundedBuffer implements Buffer {
    private static final int BUFFER_SIZE = 5;
    private int count;
    private int in;
    private int out;
    private Object[] buffer;
    public BoundedBuffer() { }
    public void insert(Object item) { }
    public Object remove() { }
}
```



```
public BoundedBuffer() {
    count = 0; in = 0; out = 0;
    buffer = new Object[BUFFER_SIZE];
}
```

```
public void insert(Object item) {
    while (count == BUFFER_SIZE);
    ++count;
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
}
```

```
public Object remove() {
    Object item;
    while (count == 0);
    --count;
    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    return item;
}
```


The Producer

```
import java.util.*;  
public class Producer implements Runnable {  
    private Buffer buffer;
```

```
    public Producer(Buffer b) {  
        buffer = b;  
    }
```

```
    public void run() {  
        Date message;  
        while (true) {  
            int sleeptime = (int) (5000 * Math.random());  
            System.out.println("Producer napping "+sleeptime+" milliseconds");  
            try {Thread.sleep(sleeptime); }  
            catch (InterruptedException e){ }  
            message = new Date();  
            System.out.println("Producer is a wake and produced " + message);  
            buffer.insert(message);  
        }  
    }
```

```
}
```

The Consumer

```
import java.util.*;
public class Consumer implements Runnable {
    private Buffer buffer;

    public Consumer(Buffer b) {
        buffer = b;
    }

    public void run() {
        Date message;
        while (true) {
            int sleeptime = (int) (5000 * Math.random());
            System.out.println("Consumer napping "+sleeptime+" milliseconds");
            try {sleep(sleeptime);}
            catch (InterruptedException e){}
            System.out.println("Consumer is awake and wants to consume.");
            message = (Date) buffer.remove();
        }
    }
}
```

The Factory Class

```
public class Factory{
```

```
    public Factory(){  
        Buffer buffer = new BoundedBuffer();  
        Thread producerThread = new Thread(new Producer(buffer));  
        Thread consumerThread = new Thread(new Consumer(buffer));  
  
        producerThread.start();  
        consumerThread.start();  
    }
```

```
    public static void main(String args[]) {  
        new Factory();  
    }
```

```
}
```


The Problem

- ❖ Concurrent access to shared data may result in data inconsistency
- ❖ Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- ❖ Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **count** that keeps track of the number of full buffers. Initially, count is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

Threads accessing shared data

```
while (count == BUFFER_SIZE)
    ; // do nothing


// add an item to the buffer
++count;
buffer[in] = item;
in = (in + 1) % BUFFER_SIZE;
```



Body of the insert method
invoked by the Producer

```
while (count == 0)
    ; // do nothing

// remove an item from the buffer
--count;
item = buffer[out];
out = (out + 1) % BUFFER_SIZE;
```



Body of the remove
method invoked
by the Consumer

A scenario demonstrating the problem

- ❖ `count++` could be implemented as

```
register1 = count  
register1 = register1 + 1  
count = register1
```

- ❖ `count--` could be implemented as

```
register2 = count  
register2 = register2 - 1  
count = register2
```

- ❖ Consider this execution interleaving with “count = 5” initially:

```
S0: producer execute register1 = count {register1 = 5}  
S1: producer execute register1 = register1 + 1 {register1 = 6}  
S2: consumer execute register2 = count {register2 = 5}  
S3: consumer execute register2 = register2 - 1 {register2 = 4}  
S4: producer execute count = register1 {count = 6}  
S5: consumer execute count = register2 {count = 4}
```

Synchronization Terminology

- ❖ Concurrent accesses to same shared variable, where at least one access is a write:
 - Order of accesses may change result of program
 - May cause irregular errors, very hard to debug
- ❖ Definition of Race Condition: the situation when there is concurrent access to shared data and the final outcome depends upon order of execution.
- ❖ Definition of Critical Section: Section of code where shared data is accessed.
- ❖ Definition of Synchronization: Mechanism that allows the programmer to control the relative order in which operations occur in different threads or processes. (Coordinate Access of shared data)

Cooperative Processes Interaction Structure

```
while (true) {
```

entry section

critical section

exit section

remainder section

```
}
```

...

```
while (true) {
```

entry section

critical section

exit section

remainder section

```
}
```

Process/thread 1

...

Process/thread n

- ❑ **Entry Section** - Code that requests permission to enter its critical section.
- ❑ **Exit Section** - Code that runs after exiting the critical section


The Concept of a Lock

- ❖ **Definition of a Lock:** LockEntity can be held by only one thread at a time granting a type of permission to do something. LockEntity has an operation to **acquire/hold/lock** and another operation to **release/give-away/unlock**.
- ❖ **Properties:**
 - A type of synchronization
 - Used to enforce mutual exclusion
 - Thread can acquire / release locks
 - Thread will wait to acquire lock (stop execution)
 - If lock held by another thread



Java Synchronization

- ❖ Java provides synchronization at the language-level.
- ❖ Each Java object has an associated lock.
 - This lock is acquired by invoking a **synchronized** method.
 - This lock is released when exiting the synchronized method.
- ❖ Threads waiting to acquire the object lock are placed in the **entry set** for the object lock.



Each object has an associated **entry set**.

Synchronized Methods In Java

```
Public synchronized void insert(Object item){  
    ..... // body of the method goes here  
}
```




Short hand notation for

```
Public void insert(Object item){  
    synchronized(this){  
        ..... // body of the method goes here  
    }  
}
```

Locks in Java

❖ Properties:

- No other thread can get lock in synchronized block
- Locked block of code  critical section

❖ Lock is released when block terminates:

- End of block reached
- Exit block due to ***return, continue, break***
- Exception thrown

Example: Producer/Consumer Problem

Synchronized insert() and remove() methods

```
public synchronized void insert(Object item) {
    while (count == BUFFER_SIZE)
        Thread.yield();

    ++count;
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
}

public synchronized Object remove() {
    Object item;

    while (count == 0)
        Thread.yield();

    --count;
    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    return item;
}
```

Java Synchronization wait/notify()

❖ When a thread invokes *wait()*:


1. The thread releases the object lock;
2. The state of the thread is set to Blocked;
3. The thread is placed in the **wait set** for the object.

❖ When a thread invokes *notify()*:

1. An arbitrary thread T from the wait set is selected;
2. T is moved from the wait to the entry set;
3. The state of T is set to Runnable.

❖ When a thread invokes *notifyAll()*:

- selects all threads in the wait set and moves them to the entry set.



Java Synchronization wait/notify()

❖ Synchronization rules in Java:

- A thread that owns the lock for an object can enter another synchronized method (or block) for the same block. This is known recursive or reentrant lock.
- A thread can nest synchronized method invocations for different objects. Thus, a thread can simultaneously own the lock of several different objects.
- If a method is not declared synchronized, then it can be invoked regardless of lock ownership, even when another synchronized method for the same object is executing.
- If the wait set of an object is empty, then the call for notify() or notifyall() has no effect.
- wait(), notify(), and notifyall() may only be invoked from synchronized methods or blocks; otherwise, an IllegalMonitorStateException is thrown.

Producer/Consumer Problem

```
public synchronized void insert(Object item) {
    while (count == BUFFER_SIZE) {
        try {
            wait();
        }
        catch (InterruptedException e) { }
    }

    ++count;
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;

    notify();
}

public synchronized Object remove() {
    Object item;

    while (count == 0) {
        try {
            wait();
        }
        catch (InterruptedException e) { }
    }

    --count;
    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    notify();

    return item;
}
```


Java Synchronization explicit lock

❖ Using

- **ReentrantLock** class of **java.util.concurrent.locks** which is an implementation of the interface **java.util.concurrent.locks.Lock**
Common methods are:

- **lock(), unlock(), newCondition()**

&

- **Condition** class of **java.util.concurrent.locks** Common methods are:
 - **await(), signal(), signalAll()**

- ❖ In this technique a thread can wait using **await()** for a specific condition (which releases the lock) and can get to wake up when it is signaled for that condition using **signal()/signalAll()** if it was waiting for that condition.

Reentrant Locks

- **import**
 - `java.util.concurrent.locks.ReentrantLock;`
 - `Java.util.concurrent.locks.Lock;`
 - `Java.util.concurrent.locks.Condition;`
- **Safe usage**

```
Lock myLock = new ReentrantLock();
Condition cond1 = myLock.newCondition();
...
Condition cond2 = myLock.newCondition();
...
myLock.lock();
try{
    //critical section
    //cond1.signal();
    //cond2.signal();
    //cond1.signalAll();
    //cond2.signalAll();
}
finally{
    myLock.unlock();
}
```

The Concept of a Semaphore

- ❖ A Semaphore is a synchronization tool that does not require busy waiting
- ❖ Semaphore S – integer variable
- ❖ Two standard operations modify S :
 - `acquire()` and `release()`
 - Originally called `P()` and `V()`
- ❖ Less complicated
- ❖ Can only be accessed via two indivisible (atomic) operations

In Dutch *proberen*
(test)

In Dutch *verhogen*
(increment)

```
acquire() {  
    while value <= 0  
        ; // no-op  
    value--;  
}  
  
release() {  
    value++;  
}
```

Semaphore as General Synchronization Tool

- ❖ **Counting** semaphore – integer value can range over an unrestricted domain
- ❖ **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
 - Also known as **mutex locks**

Usage of Semaphores

```
Semaphore S = new Semaphore();  
  
S.acquire();  
  
    // critical section  
  
S.release();  
  
    // remainder section
```

Semaphore Implementation

- ❖ Must guarantee that no two processes can execute `acquire ()` and `release ()` on the same semaphore at the same time.

Busy wait: while a process is in its critical section, any other process trying to enter its critical section must loop continuously in the entry code.

Spin Lock: is a semaphore with busy wait, process spins while Waiting for the lock.

```
acquire() {  
    while value <= 0  
        ; // no-op  
    value--;  
}  
  
release() {  
    value++;  
}
```



In a single CPU multiprocessing System where locks are expected to be held for long time – critical section is long -, busy wait wastes CPU cycles that some other process might be able to use productively.



In a multiprocessor System where locks are expected to be held for short time – critical section is short -, busy wait is useful requiring no context switching.



But in most applications, locks are expected to be held for long time – critical section is long – So, this implementation is not good.

Semaphore Implementation with no Busy waiting

- ❖ With each semaphore there is an associated waiting queue. Each entry in a waiting queue has two data items:



- value (of type integer)
- pointer to next record in the list

- ❖ Two operations:



- **block** – place the process invoking the operation on the appropriate waiting queue.
- **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

Semaphore Implementation with no Busy waiting (Cont.)



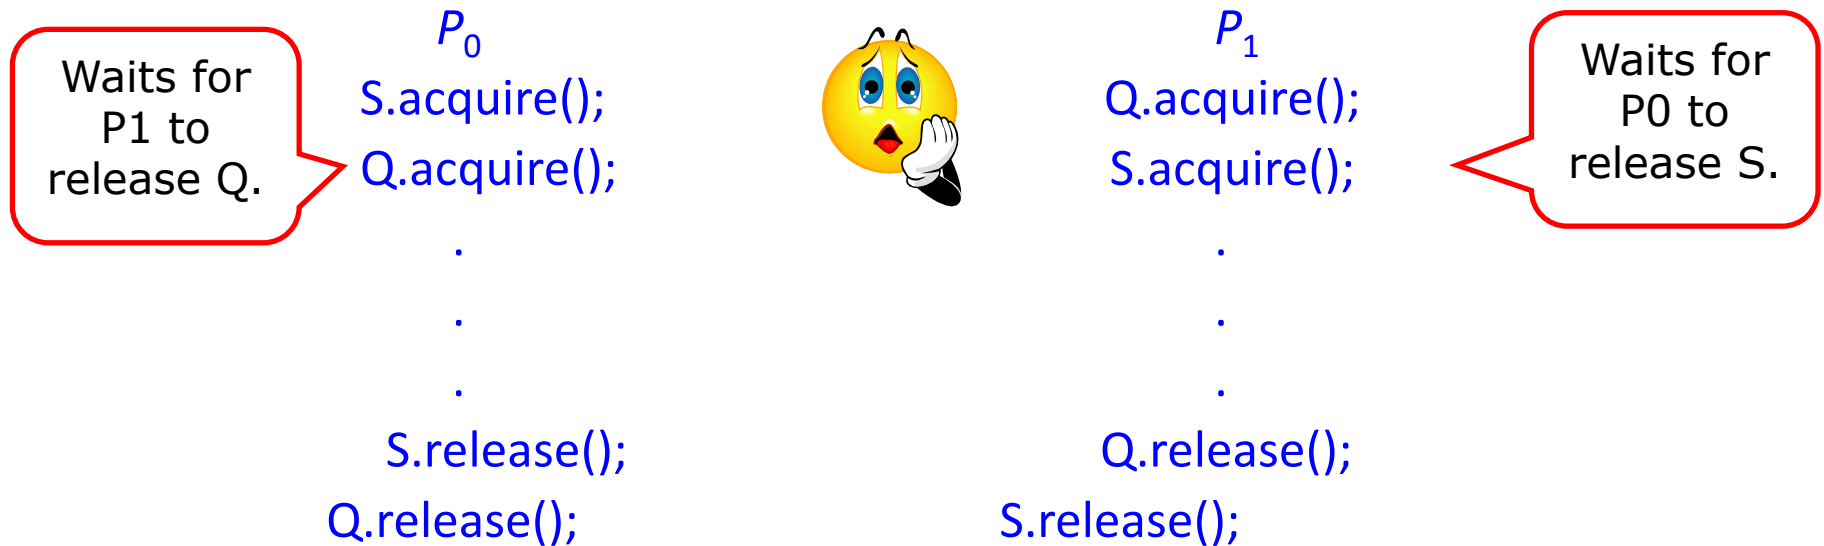
```
acquire(){
    value--;
    if (value < 0) {
        add this process to list
        block;
    }
}
```

```
release(){
    value++;
    if (value <= 0) {
        remove a process P from list
        wakeup(P);
    }
}
```

$meaning(Value) = \begin{cases} \text{Number of available resources,} & value \geq 0 \\ |value| \text{ is the number of waiting processes,} & value \leq 0 \end{cases}$

Deadlock and Starvation

- ❖ **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- ❖ Let **S** and **Q** be two semaphores initialized to 1.



- ❖ **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Semaphores

- ❖ To create a semaphore call the constructor
 - **Semaphore(int value)**, where value is allowed to be negative.
- ❖ To acquire a permit from this semaphore, blocking until one is available, or the thread is interrupted
 - **void acquire()**
- ❖ To acquire the given number of permits from this semaphore, blocking until all are available, or the thread is interrupted
 - **void acquire(int permits)**
- ❖ To return the current number of permits available in this semaphore
 - **int availablePermits()**
- ❖ To release a permit, returning it to the semaphore
 - **void release()**
- ❖ To release the given number of permits, returning them to the semaphore
 - **void release(int permits)**

Semaphores


❖ How to?

- `import java.util.concurrent.Semaphore;`
- Save usage:

```
Semaphore sem = new Semaphore(1);  
//.....  
try{  
    sem.acquire();  
    //critical section  
}  
  
    finally{  
        sem.release();  
}
```

Dining Philosophers Problem


- ❑ 5 philosophers with 5 chopsticks placed between them to eat bowl of rice (data set).
- ❑ To eat requires two chopsticks.
- ❑ Philosophers alternate between thinking and eating.
- ❑ OS examples: simultaneous use of multiple resources.



The Dining Philosophers Problem

- ❖ Shared data
 - Bowl of rice (data set)
 - Semaphore `chopStick [5]` initialized to 1

```
while (true) {  
    // get left chopstick  
    chopStick[i].acquire();  
    // get right chopstick  
    chopStick[(i + 1) % 5].acquire();  
  
    eating();  
  
    // return left chopstick  
    chopStick[i].release();  
    // return right chopstick  
    chopStick[(i + 1) % 5].release();  
  
    thinking();  
}
```



**Structure of
Philosopher *i***



The Producer/Consumer Problem

- Semaphore **mutex**
initialized to the value 1
- Semaphore **full**
initialized to the value 0.
- Semaphore **empty**
initialized to the value N.

```
public class BoundedBuffer implements Buffer
{
    private static final int BUFFER_SIZE = 5;
    private Object[] buffer;
    private int in, out;
    {
        private Semaphore mutex;
        private Semaphore empty;
        private Semaphore full;
    }

    public BoundedBuffer() {
        // buffer is initially empty
        in = 0;
        out = 0;
        buffer = new Object[BUFFER_SIZE];

        {
            mutex = new Semaphore(1);
            empty = new Semaphore(BUFFER_SIZE);
            full = new Semaphore(0);
        }

        public void insert(Object item) {
            // Figure 6.9
        }

        public Object remove() {
            // Figure 6.10
        }
    }
}
```

The Producer/Consumer Problem

```
public void insert(Object item)
    empty.acquire();
    mutex.acquire();

    // add an item to the buffer
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;

    mutex.release();
    full.release();
}
```

```
public Object remove() {
    full.acquire();
    mutex.acquire();

    // remove an item from the buffer
    Object item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    mutex.release();
    empty.release();

    return item;
}
```

Atomic variables

- ❖ An AtomicBoolean is a boolean value that may be updated atomically.
- ❖ It is used in applications such as atomically updated flags, and cannot be used as a replacement for a Boolean.
- ❖ How to?
 - **import java.util.concurrent.atomic.AtomicBoolean;**
 - To create a new AtomicBoolean with a given initial value:
 - **AtomicBoolean(boolean initialValue)**
 - To atomically set the value to the given updated value if the current value equals the expected value:
 - **boolean compareAndSet(boolean expect, boolean update)**
 - To set to the given value and returns the previous value:
 - **boolean getAndSet(boolean newValue)**
 - To return the current value:
 - **boolean get()**
 - To unconditionally set to the given value:
 - **set(boolean newValue)**

Atomic variables

❖ Other atomic variables

- **`java.util.concurrent.atomic.AtomicInteger`**
- **`java.util.concurrent.atomic.AtomicLong`**
- **`java.util.concurrent.atomic.AtomicIntegerArray`**
- **`java.util.concurrent.atomic.AtomicLongArray`**